MedAccess AI

Pitch Speaker Notes

18 slides · ~7 minutes of talk + live demo · each entry: what to **say**, and the one **beat** that must land.

Delivery

- Total: about 7 minutes of talk, then the live demo. If your slot is 5 minutes, compress slides 10-11 and 15-16 to one sentence each — never cut slides 2, 6, or 13.
- Slide 6 is the money slide. Slow down, point at the five steps, and say the beat line word for word.
- Let the stats breathe: after '4.5 billion' and after 'eighty percent', pause for one beat.
- Keys during the talk: arrows or space to advance, F for fullscreen. The deck auto-hides its hints.

01 Healthcare for the next billion patients

~25s

Hi, we're MedAccess AI. When you think of clinical AI, you probably picture tools for American hospitals — fast internet, English, an EHR behind everything. We built the opposite: an AI doctor copilot for the clinics that have the least — rural, multilingual, often one nurse for an entire district. And our agent doesn't stop at advice. It books the patient into a real clinic. That's the whole pitch in one line: healthcare that closes the loop, for the next billion patients.

Land: We close the loop — advice becomes an appointment.

---

02 The incumbents serve the already well-served

~25s

Every serious clinical-AI company today — OpenEvidence, Glass Health, Hippocratic — is racing for the same customer: the well-served Western clinician. Meanwhile the WHO counts 4.5 billion people without full access to essential health services, against a projected shortfall of more than ten million health workers by 2030. And the number of major clinical-AI products built for that population? Zero. That gap is our wedge.

Land: 4.5B people, 10M missing clinicians, zero products. That's the wedge.

---

03 Three pain points define a rural clinic's day

~25s

On the ground it comes down to three problems. Too few clinicians — minutes per patient, no second opinion within hours of travel. Language — the provider and the patient often can't run a proper diagnostic interview in a shared language. And the cruelest one: even when a patient gets a good read, there's no path to care. No referral, no booking, no follow-up. Advice without an appointment changes nothing.

Land: The unsolved part isn't the diagnosis — it's the path to care.

---

04 We aren't competing for the US specialist

~20s

Here's the competitive picture in one table. Every incumbent: English-first, text-first, no medical images, online-only, built for the US or EU. We're multilingual — seventeen-plus languages — voice-first, multimodal, and installable as a PWA on any phone. We're not out-competing these companies at their own game; we're serving the patients they structurally cannot reach.

Land: Different patients, not a better version of the same product.

---

05 Meet MA Agent ~25s

So meet MA Agent — the colleague a lone clinician never had. It runs in any browser, installs on any phone, speaks the patient's language, listens to voice, reads X-rays and skin photos — and books appointments. One deliberate design choice: patients always talk to 'MA Agent', a named, consistent identity — never 'a chatbot from company X'. Under the hood it's four pieces: a patient PWA, a clinic portal, one shared API, and a Python sidecar running specialist vision models. And to be clear: it's decision support. It never replaces a licensed clinician.

Land: A named, trusted colleague — not a chatbot.

---

06 Most clinical AI stops at the read. We close the loop. ~30s

This is the slide I want you to remember. Most clinical AI stops at the read — we close the loop. Watch the five steps: the patient describes symptoms in any language, by text, voice, or photo. After a few turns the agent forms a clinical snapshot — specialty plus urgency. Then the differentiator: it offers a real open slot with the right doctor. The patient says yes. An urgency-badged referral lands in that clinic's queue with full context. The doctor confirms in one tap. And if no enrolled clinic is nearby, we fall back to public maps with one-tap navigation — the patient always has somewhere to go.

Land: The read isn't the product. The booking is.

---

07 Booking happens inside the conversation ~25s

How does booking actually work? Inside the conversation — no forms, no buttons. The agent is handed the doctor's real availability, so it can only offer slots that actually exist — it cannot invent a time. The moment the patient agrees, a hidden marker in the model's reply silently calls our booking API. What you see on this screen is the real product experience: three messages, and the appointment exists.

Land: Three chat bubbles = a confirmed appointment. No forms.

---

08 One copilot, the whole clinical journey ~20s

Zooming out: one copilot, seven modules across the journey. Patient side — the MA Agent chat and Find Care with the maps fallback. Clinic side — the urgency-badged patients queue, a structured interview module, symptom analysis with calibrated differentials, report reading, and Manchester-style triage. Every module feeds the same loop: patient to clinic.

Land: Not a feature — a full journey, both sides of the visit.

---

09 Generalist LLM plus specialist medical models ~30s

On medical images we made a bet the research supports. Generic vision LLMs plateau around seventy to ninety percent on condition-specific reads. Purpose-built models reach ninety-two plus. So we run specialist models locally: chest X-ray across eighteen pathologies, skin lesions, diabetic retinopathy — all three live in this demo build — with the malaria smear model queued. The important part: the language model never diagnoses from pixels. The specialist model reads; the LLM explains in plain language; and any disagreement is flagged, never silently overridden. These five diseases drive roughly eighty percent of visits in our target clinics.

Land: Specialists detect, the LLM only explains — disagreements are flagged.

---

10 Clean monorepo. One contract. Graceful degradation. ~20s

Architecture in one breath: two React PWAs, one Express API, and a single Zod schema contract shared end to end — zero duplicated types across three apps. MongoDB is fire-and-forget, retrieval is BM25 over a curated clinical corpus, and the Python sidecar serves the vision models. Everything degrades gracefully: kill the database, the sidecar, or voice — the app keeps working.

Land: One schema contract; nothing is a single point of failure.

---

11 Every decision earns its place ~20s

A few engineering decisions that earn their place. An LLM gateway means one environment variable swaps the model — judges can A/B models live on the same prompt. A cheap-stack default means the full demo costs effectively zero to run. BM25 instead of a vector database, because at our corpus size keyword retrieval matches embeddings with no extra infrastructure. A PWA instead of native — five-second install on an entry-level Android. And no accounts: a phone number is the only identity a patient needs to book.

Land: Pragmatic choices, each defensible under questioning.

---

12 Trustworthy by construction ~25s

A medical copilot has to earn trust by construction, not by disclaimer. Four guardrails. A hard triage floor: chest pain, stroke-like deficits, severe breathing difficulty can never be triaged below orange — that's enforced in the prompt contract. A held identity: the system never reveals which model or provider is underneath. Calibrated uncertainty: 'most likely, possible, unlikely' — never a fake ninety-five percent, never an invented dosage or citation. And privacy by default: sessions expire, chat history lives on the patient's device.

Land: Safety is enforced in the system's constitution, not promised in fine print.

---

13 v1.0 demo: the headline loop works end to end ~20s

We grade ourselves honestly. Done: the full patient-to-clinic loop — chat, snapshot, conversational booking, clinic queue, confirmation. Done: streaming chat with RAG citations, push-to-talk voice, three live specialist image models, the maps fallback, and clean typechecks across every workspace. In progress: visual polish. Queued: the malaria weights and a skin-model upgrade. Nothing on this slide is aspirational — you'll watch it run in a minute.

Land: Every 'DONE' on this slide is about to be demonstrated live.

---

14 If we reach even a fraction, the math is enormous ~25s

Why this matters at scale. In endemic regions, malaria alone is around forty percent of pediatric outpatient visits. Pneumonia is fifteen percent of under-five visits in rural South Asia. Our five Pareto diseases cover roughly eighty percent of what walks into a target clinic — and the hardware to serve an entire clinic is one two-hundred-dollar laptop on local WiFi. Every screening that catches a red flag early, or routes a patient to the right doctor today instead of next month, is a life trajectory changed.

Land: Five diseases ~ 80% of visits; $200 serves a whole clinic.

---

15 A built-in network effect

~20s

Distribution compounds. Patients arrive through a shared link or QR code — zero install, zero friction. The maps fallback shows every nearby clinic — which is also our recruitment list. Every clinic that enrolls gets free, pre-triaged referrals it would never have reached, which makes in-app booking better, which pulls more patients. We sit in the middle as the routing layer.

Land: The fallback map is also the sales pipeline.

---

16 A credible path from demo to deployment

~20s

Where this goes. Today: the full closed loop you're about to see — v1.0. Next weeks: authentication, the pharmacy referral loop, the malaria model, and SMS confirmations. Months out: a production-scale RAG corpus, clinician-signed summaries, reverse patient matching. With funding: compliance work, native mobile, offline on-device speech, and EHR write-back. Each step is engineering, not research risk.

Land: Everything ahead is execution, not invention.

---

17 Five engineers, clear lanes

~15s

The team: five engineers with clear lanes — architecture, the ML models, product UI, ops and QA, and docs — under one discipline: schemas first, strict types, conventional commits, and no refactors of working code two weeks before a deadline. Even our AI coding agents are role-scoped: every session starts by asking which engineer it's helping.

Land: Small team, unusual discipline.

---

18 Let's watch it close the loop

~15s + demo

That's the story. Now let's watch it close the loop — live. On the left screen, a patient describes symptoms on a phone. On the right, the clinic portal. You'll see the conversation become a booking, and the booking become a confirmed appointment in the doctor's queue. And if you'd like to try it yourself, scan the QR on the demo screen — it runs on your phone, right now. Thank you.

Land: End on the live loop; invite them to scan the QR.
